Tentative NumPy Tutorial

Please do not hesitate to click the *edit* button. You will need to create a User Account first.

Contents

Prerequisites

Before reading this tutorial you should know a bit of Python. If you would like to refresh your memory, take a look at the Python tutorial.

You also need to have some software installed on your computer. As a bare minimum you would need:

- Python
- NumPy

These you may find useful:

- ipython is an enhanced interactive Python shell which is very convenient for exploring NumPy's features
- matplotlib will enable you to plot graphics
- SciPy provides a lot of scientific routines working on the top of NumPy

The Basics

NumPy's main object is the homogeneous multidimensional array. This is a table of elements (usually numbers), all of the same type, indexed by a tuple of positive integers.

By 'multidimensional', we mean that arrays can have several *dimensions* or *axes*. Because the word *dimension* is ambiguous, we use **axis** instead. The number of axes will often be called *rank*.

For example, the coordinates of a point in 3D space `[1, 2, 1]` is an array of rank 1 ---it has one axis. That axis has a length of 3. As another example, the array

```
[[ 1., 0., 0.],
 [ 0., 1., 2.]]
```

is an array of rank 2 (it is 2-dimensional). The first dimension (axis) has a length of 2, the second dimension has a length of 3. For more details see the Numpy Glossary.

The multidimensional array class is called `ndarray`. Note that this is not the same as the Standard Python Library class `array`, which is only for one-dimensional arrays. The more important attributes of an `ndarray` object are:

ndarray.ndim
> the number of axes (dimensions) of the array. In the Python world, the number of dimensions is often referred to as *rank*.

ndarray.shape
> the dimensions of the array. This is a tuple of integers indicating the size of the array in each dimension. For a matrix with *n* rows and *m* columns, `shape` will be `(n,m)`. The length of the `shape` tuple is therefore the rank, or number of dimensions, `ndim`.

ndarray.size
> the total number of elements of the array. This is equal to the product of the elements of `shape`.

ndarray.dtype
> an object describing the type of the elements in the array. One can create or specify dtype's using standard Python types. NumPy provides a bunch of them, for example: `bool_`, `character`, `int_`, `int8`, `int16`, `int32`, `int64`, `float_`, `float8`, `float16`, `float32`, `float64`, `complex_`, `complex64`, `object_`.

ndarray.itemsize

the size in bytes of each element of the array. For example, an array of elements of type `float64` has `itemsize` 8 (=64/8), while one of type `complex32` has `itemsize` 4 (=32/8). It is equivalent to `ndarray.dtype.itemsize`.

ndarray.data
the buffer containing the actual elements of the array. Normally, we won't need to use this attribute because we will access to the elements in an array using indexing facilities.

An example

We define the following array:

```
>>> from numpy  import *
>>> a = arange(10).reshape(2,5)
>>> a
array([[0, 1, 2, 3, 4],
       [5, 6, 7, 8, 9]])
```

We have just created an array object with a label *a* attached to it. The array *a* has several attributes --or properties. In Python, attributes of a specific object are denoted `name_object.attribute`. In our case:

- `a.shape` is (2,5)
- `a.ndim` is 2 (which is the length of `a.shape`)
- `a.size` is 10
- `a.dtype.name` is `int32`
- `a.itemsize` is 4, which means that an `int32` takes 4 bytes in memory.

You can check all these attributes just by typing them in the interactive session:

```
>>> a.shape
(2, 5)
>>> a.dtype.name
'int32'
```

And so on.

Array Creation

There are many ways to create arrays. For example, you can create an array from a regular Python list or tuple using the `array` function.

```
>>> from numpy  import *
>>> a = array( [2,3,4] )
>>> a
array([2, 3, 4])
>>> type(a)                                    # a is an object
of the ndarray class
<type 'numpy.ndarray'>
```

`array` transforms sequences of sequences into two dimensional arrays, and it transforms sequences of sequences of sequences into three dimensional arrays, and so on. The type of the resulting array is deduced from the type of the elements in the sequences.

```
>>> b = array( [ (1.5,2,3), (4,5,6) ] )          # this will be an
array of floats
>>> b
array([[ 1.5,  2. ,  3. ],
       [ 4. ,  5. ,  6. ]])
```

Once we have an array we can take a look at its attributes:

```
>>> b.ndim                                       # number of
dimensions
2
>>> b.shape                                      # the dimensions
(2, 3)
>>> b.dtype                                      # the type (8 byte
floats)
dtype('float64')
>>> b.itemsize                                   # the size the
type
8
```

The type of the array can also be explicitly specified at creation time:

```
>>> c = array( [ [1,2], [3,4] ], dtype=complex )
>>> c
array([[ 1.+0.j,  2.+0.j],
       [ 3.+0.j,  4.+0.j]])
```

A frequent error consists in calling `array` with a multiple numeric arguments, rather than providing a single list of numbers as an argument.

```
>>> a = array(1,2,3,4)    # WRONG
#!python numbers=disable
>>> a = array([1,2,3,4])  # RIGHT
```

The function `array` is not the only one that creates arrays. Usually the elements of the array are not known from the beginning, and a placeholder array is needed. There are some functions to create arrays with some initial content. By default, the type of the created array is `float64`.

The function `zeros` creates an array full of zeros, and the function `ones` creates an array full of ones.

```
>>> zeros( (3,4) )                               # the parameter
specifies the shape
array([[0.,  0.,  0.,  0.],
       [0.,  0.,  0.,  0.],
       [0.,  0.,  0.,  0.]])
```

```
>>> ones( (2,3,4), dtype=int16 )                    # dtype can also
be specified
array([[[ 1, 1, 1, 1],
        [ 1, 1, 1, 1],
        [ 1, 1, 1, 1]],
       [[ 1, 1, 1, 1],
        [ 1, 1, 1, 1],
        [ 1, 1, 1, 1]]], dtype=int16)
```

The function `empty` creates an array without filling it in. Then the initial content is random and it depends on the state of the memory.

```
>>> empty( (2,3) )
array([[  3.73603959e-262,   6.02658058e-154,   6.55490914e-260],
       [  5.30498948e-313,   3.14673309e-307,   1.00000000e+000]])
>>> empty( (2,3) )                                  # the content may
change in different invocations
array([[  3.14678735e-307,   6.02658058e-154,   6.55490914e-260],
       [  5.30498948e-313,   3.73603967e-262,   8.70018275e-313]])
```

To create sequences of numbers, NumPy provides a function analogous to `range` that returns arrays instead of lists

```
>>> arange( 10, 30, 5 )
array([10, 15, 20, 25])
>>> arange( 0, 2, 0.3 )                          # it accepts float
arguments
array([ 0. ,  0.3,  0.6,  0.9,  1.2,  1.5,  1.8])
```

Using `arange` with floating point arguments, it is generally not possible to predict the number of elements obtained (because of the floating point precision). For this reason, it is usually better to use the function `linspace` that receives as an argument the number of elements that we want, instead of the step:

```
>>> linspace( 0, 2, 9 )                          # 9 numbers from 0 to 2
array([ 0.  ,  0.25,  0.5 ,  0.75,  1.  ,  1.25,  1.5 ,  1.75,  2.
])
>>> x = linspace( 0, 2*pi, 100 )                 # useful to evaluate
function at lots of points
>>> f = sin(x)
```

See also

> array, zeros, zeros like, ones, ones like, empty, empty like, arange, linspace, rand, randn, fromfunction, fromfile

Printing Arrays

When you print an array, NumPy displays it in a similar way to nested lists, but with the following layout:

- the last axis is printed from left to right,
- the last but one, from top to bottom,
- and the rest, also from top to bottom, separating each slice by an empty line.

One dimensional arrays are then printed as rows, bidimensionals as matrices and tridimensionals as lists of matrices.

```
>>> a = arange(6)                        # 1d array
>>> print a
[0 1 2 3 4 5]
>>>
>>> b = arange(12).reshape(4,3)          # 2d array
>>> print b
[[ 0  1  2]
 [ 3  4  5]
 [ 6  7  8]
 [ 9 10 11]]
>>>
>>> c = arange(24).reshape(2,3,4)        # 3d array
>>> print c
[[[ 0  1  2  3]
  [ 4  5  6  7]
  [ 8  9 10 11]]
 [[12 13 14 15]
  [16 17 18 19]
  [20 21 22 23]]]
```

See below to get more details on `reshape`.

If an array is too large to be printed, NumPy automatically skips the central part of the array and only prints the corners:

```
>>> print arange(10000)
[   0    1    2 ..., 9997 9998 9999]
>>>
>>> print arange(10000).reshape(100,100)
[[   0    1    2 ...,   97   98   99]
 [ 100  101  102 ...,  197  198  199]
 [ 200  201  202 ...,  297  298  299]
 ...,
 [9700 9701 9702 ..., 9797 9798 9799]
 [9800 9801 9802 ..., 9897 9898 9899]
 [9900 9901 9902 ..., 9997 9998 9999]]
```

To disable this behaviour and force NumPy to print the entire array, you can change the printing options using `set_printoptions`.

```
>>> set_printoptions(threshold=nan)
```

Basic Operations

Arithmetic operators on arrays apply *elementwise*. A new array is created and filled with the result.

```
>>> a = array( [20,30,40,50] )
>>> b = arange( 4 )
>>> c = a-b
>>> c
array([20, 29, 38, 47])
>>> b**2
array([0, 1, 4, 9])
>>> 10*sin(a)
array([ 9.12945251, -9.88031624,  7.4511316 , -2.62374854])
>>> a<35
array([True, True, False, False], dtype=bool)
```

Unlike in many matrix languages, the product operator * operates elementwise in NumPy arrays. The matrix product can be performed using the `dot` function or creating `matrix` objects (see matrix section of this tutorial).

```
>>> A = array( [[1,1],
...             [0,1]] )
>>> B = array( [[2,0],
...             [3,4]] )
>>> A*B                         # elementwise product
array([[2, 0],
       [0, 4]])
>>> dot(A,B)                    # matrix product
array([[5, 4],
       [3, 4]])
```

It is possible to perform some operations inplace so that no new array is created.

```
>>> a = ones((2,3), dtype=int)
>>> b = random.random((2,3))
>>> a *= 3
>>> a
array([[3, 3, 3],
       [3, 3, 3]])
>>> b += a
>>> b
array([[ 3.69092703,  3.8324276 ,  3.0114541 ],
       [ 3.18679111,  3.3039349 ,  3.37600289]])
>>> a += b                                # b is converted to
integer type
>>> a
array([[6, 6, 6],
       [6, 6, 6]])
```

When operating with arrays of different types, the type of the resulting array corresponds to the more general or precise one (so called upcasting).

```
>>> a = ones(3, dtype=int32)
>>> b = linspace(0,pi,3)
>>> b.dtype.name
'float64'
>>> c = a+b
>>> c
array([ 1.        ,  2.57079633,  4.14159265])
>>> c.dtype.name
'float64'
>>> d = exp(c*1j)
>>> d
array([ 0.54030231+0.84147098j, -0.84147098+0.54030231j,
       -0.54030231-0.84147098j])
>>> d.dtype.name
'complex128'
```

Many unary operations, like computing the sum of all the elements in the array, are implemented as methods of the `ndarray` class.

```
>>> a = random.random((2,3))
>>> a
array([[ 0.6903007 ,  0.39168346,  0.16524769],
       [ 0.48819875,  0.77188505,  0.94792155]])
>>> a.sum()
3.4552372100521485
>>> a.min()
0.16524768654743593
>>> a.max()
0.9479215542670073
```

By default, these operations apply to the array as if it were a list of numbers, regardless of its shape. However, by specifying the `axis` parameter you can apply an operation along the specified axis of an array:

```
>>> b = arange(12).reshape(3,4)
>>> b
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>>
>>> b.sum(axis=0)                            # sum of each column
array([12, 15, 18, 21])
>>>
>>> b.min(axis=1)                            # min of each row
array([0, 4, 8])
>>>
>>> b.cumsum(axis=1)                         # cumulative sum
along the rows
array([[ 0,  1,  3,  6],
       [ 4,  9, 15, 22],
```

```
          [ 8, 17, 27, 38]])
```

Universal Functions

NumPy provides familiar mathematical functions, sin, cos, exp, etc. But in NumPy these functinions are called "universal functions"(`ufunc`). The reason for having such a name, within NumPy these functions operated on an array object elementwise, producing array as an output.

```
>>> B = arange(3)
>>> B
array([0, 1, 2])
>>> exp(B)
array([ 1.        ,  2.71828183,  7.3890561 ])
>>> sqrt(B)
array([ 0.        ,  1.        ,  1.41421356])
>>> C = array([2., -1., 4.])
>>> add(B, C)
array([ 2.,  0.,  6.])
```

See also

> all, alltrue, any, apply along axis, argmax, argmin, argsort, average, bincount, ceil, clip, conj, conjugate, corrcoef, cov, cross, cumprod, cumsum, diff, dot, floor, inner, inv, lexsort, max, maximum, mean, median, min, minimum, nonzero, outer, prod, re, round, sometrue, sort, std, sum, trace, transpose, var, vdot, vectorize, where

Indexing, Slicing and Iterating

One dimensional arrays can be indexed, sliced and iterated over pretty much like lists and other Python sequences.

```
>>> a = arange(10)**3
>>> a
array([  0,   1,   8,  27,  64, 125, 216, 343, 512, 729])
>>> a[2]
8
>>> a[2:5]
array([ 8, 27, 64])
>>> a[:6:2] = -1000                        # modify elements in a
>>> a[::-1]                                 # reversed a
array([ 729,  512,  343,  216,  125, -1000,   27, -1000,
1, -1000])
>>> for i in a:
...        print i**(1/3.),
...
nan 1.0 nan 3.0 nan 5.0 6.0 7.0 8.0 9.0
```

Multidimensional arrays can be with one index per axis. These indices are given in a tuple separated by commas:

```
>>> def f(x,y):
...         return 10*x+y
...
>>> b = fromfunction(f,(5,4),dtype=int)
>>> b
array([[ 0,  1,  2,  3],
       [10, 11, 12, 13],
       [20, 21, 22, 23],
       [30, 31, 32, 33],
       [40, 41, 42, 43]])
>>> b[2,3]
23
>>> b[:,1]                                # the second column of
b
array([ 1, 11, 21, 31, 41])
>>> b[1:3,:]                              # the second and third
row of b
array([[10, 11, 12, 13],
       [20, 21, 22, 23]])
```

When fewer indices are provided than the number of axes, the missing indices are supposed to be complete slices:

```
>>> b[-1]                                 # the last row.
Equivalent to b[-1,:]
array([40, 41, 42, 43])
```

`b[i]` can be read an `b[i, <as many :, as needed>]`. In NumPy, this can also be written using dots as `b[i,...]`.

The **dots** (`...`) mean: as many ':,' as needed to have a complete indexing tuple. For example, if `x` is a rank 5 array (it has 5 axes), then

- `x[1,2,...]` is equivalent to `x[1,2,:,:,:]`,
- `x[...,3]` to `x[:,:,:,:,3]` and
- `x[4,...,5,:]` to `x[4,:,:,5,:]`.

```
>>> c = array( [ [[  0,  1,  2],          # a 3d array (two
stacked 2d arrays)
...               [ 10, 12, 13]],
...
...              [[100,101,102],
...               [110,112,113]] ] )
>>> c.shape
(2, 2, 3)
>>> c[1,...]                              # same as c[1,:,:]
or c[1]
array([[100, 101, 102],
       [110, 112, 113]])
```

```
>>> c[...,2]                              # same as c[:,:,2]
array([[  2,  13],
       [102, 113]])
```

Iterating over multidimensional arrays is done with respect to the first axis:

```
>>> for row in b:
...         print row
...
[0 1 2 3]
[10 11 12 13]
[20 21 22 23]
[30 31 32 33]
[40 41 42 43]
```

However, if one wants to perform something for each element in the array, one can use the `flat` attribute which is an iterator over all the elements of the array:

```
>>> for element in b.flat:
...         print element,
...
0 1 2 3 10 11 12 13 20 21 22 23 30 31 32 33 40 41 42 43
```

See also
    [], ..., newaxis, ndenumerate, indices, index exp

Shape Manipulation

Changing the shape of an array

An array has a shape, given by the number of elements along each axis:

```
>>> a = floor(10*random.random((3,4)))
>>> a
array([[ 7.,  5.,  9.,  3.],
       [ 7.,  2.,  7.,  8.],
       [ 6.,  8.,  3.,  2.]])
>>> a.shape
(3, 4)
```

The shape of an array can be changed with various commands:

```
>>> a.ravel() # flatten the array
array([ 7.,  5.,  9.,  3.,  7.,  2.,  7.,  8.,  6.,  8.,  3.,
2.])
>>> a.shape = (6, 2)
>>> a.transpose()
array([[ 7.,  9.,  7.,  7.,  6.,  3.],
       [ 5.,  3.,  2.,  8.,  8.,  2.]])
```

The order of the elements in the array resulting from ravel() is normally "C-style", that is, the rightmost index "changes the fastest", so that the element after a[0,0] is a[0,1]. If the array is reshaped to some other shape, again the array is treated as "C-style". Numpy normally creates arrays stored in this order, so ravel() will usually not need to copy its argument, but if the array was made by taking slices of another array or created with unusual options, it may need to be copied. The functions ravel() and reshape() can also be told (using an optional argument) to use FORTRAN-style arrays, in which the leftmost index changes the fastest.

The reshape function returns its argument with a modified shape, whereas the resize method modifies the array itself:

```
>>> a
array([[ 7.,   5.],
       [ 9.,   3.],
       [ 7.,   2.],
       [ 7.,   8.],
       [ 6.,   8.],
       [ 3.,   2.]])
>>> a.resize((2,6))
>>> a
array([[ 7.,   5.,   9.,   3.,   7.,   2.],
       [ 7.,   8.,   6.,   8.,   3.,   2.]])
```

If in a reshaping operation a dimension is given as -1, it is automatically calculated to correspond to the other dimensions:

```
>>> a.reshape(3,-1)
array([[ 7.,   5.,   9.,   3.],
       [ 7.,   2.,   7.,   8.],
       [ 6.,   8.,   3.,   2.]])
```

See also:: shape example, reshape example, resize example, ravel example

Stacking together different arrays

Several arrays can be stacked together, along different axes:

```
>>> a = floor(10*random.random((2,2)))
>>> a
array([[ 1.,   1.],
       [ 5.,   8.]])
>>> b = floor(10*random.random((2,2)))
>>> b
array([[ 3.,   3.],
       [ 6.,   0.]])
>>> vstack((a,b))
array([[ 1.,   1.],
       [ 5.,   8.],
       [ 3.,   3.],
       [ 6.,   0.]])
```

```
>>> hstack((a,b))
array([[ 1.,   1.,   3.,   3.],
       [ 5.,   8.,   6.,   0.]])
```

The function column stack stacks 1D arrays as columns into a 2D array. It is equivalent to vstack only for 1D arrays:

```
>>> column_stack((a,b))    # With 2D arrays
array([[ 1.,   1.,   3.,   3.],
       [ 5.,   8.,   6.,   0.]])
>>> a=array([4.,2.])
>>> b=array([2.,8.])
>>> a[:,newaxis]  # This allows to have a 2D columns vector
array([[ 4.],
       [ 2.]])
>>> column_stack((a[:,newaxis],b[:,newaxis]))
array([[ 4.,   2.],
       [ 2.,   8.]])
>>> vstack((a[:,newaxis],b[:,newaxis])) # The behavior of vstack
is different
array([[ 4.],
       [ 2.],
       [ 2.],
       [ 8.]])
```

The function row stack, on the other hand, stacks 1D arrays as rows into a 2D array.

For arrays of with more than two dimensions, hstack stacks along their first axes, vstack stacks along their last axes, and concatenate allows for an optional arguments giving the number of the axes along which the concatenation should happen.

Note

In complex cases, r_[] and c_[] are useful for creating arrays by stacking numbers along one axis. They allow the use of range literals (":") :

```
>>> r_[1:4,0,4]
array([1, 2, 3, 0, 4])
```

When used with arrays as arguments, r[] and c[] are similar to vstack and hstack in their default behavior, but allow for a optional arguments giving the number of the axis along which to concatenate.

See also: hstack example, vstack exammple, column_stack example, row_stack example, concatenate example, c_ example, r_ example

Splitting one array into several smaller ones

Using hsplit you can split an array along its horizontal axis, either by specifying the number of equally shaped arrays to return, or by specifying the columns after which the division should occur:

```
>>> a = floor(10*random.random((2,12)))
```

```
>>> a = floor(10*random.random((2,12)))
>>> a
array([[ 8.,   8.,   3.,   9.,   0.,   4.,   3.,   0.,   0.,   6.,   4.,
         4.],
       [ 0.,   3.,   2.,   9.,   6.,   0.,   4.,   5.,   7.,   5.,   1.,
         4.]])
>>> hsplit(a,3)   # Split a into 3
[array([[ 8.,   8.,   3.,   9.],
        [ 0.,   3.,   2.,   9.]]), array([[ 0.,   4.,   3.,   0.],
        [ 6.,   0.,   4.,   5.]]), array([[ 0.,   6.,   4.,   4.],
        [ 7.,   5.,   1.,   4.]])]
>>> hsplit(a,(3,4))   # Split a after the third and the fourth
column
[array([[ 8.,   8.,   3.],
        [ 0.,   3.,   2.]]), array([[ 9.],
        [ 9.]]), array([[ 0.,   4.,   3.,   0.,   0.,   6.,   4.,   4.],
        [ 6.,   0.,   4.,   5.,   7.,   5.,   1.,   4.]])]
```

vsplit splits along the vertical axis, and array split allows to specify along which axis to split.

Copies and Views

When operating and manipulating arrays, their data is sometimes copied into a new array and sometimes not. This is often a source of confusion for beginners. There are three cases:

No Copy at All

Simple assignments don't make any copy of array objects nor of their data.

```
>>> a = arange(12)
>>> b = a                              # no new object is
created
>>> b is a                             # a and b are two names
for the same ndarray object
True
>>> b.shape = 3,4                      # changes the shape of a
>>> a.shape
(3, 4)
```

Python passes mutable objects as references, so function calls make no copy.

```
>>> def f(x):
...     print id(x)
...
>>> id(a)                              # id is a unique
identificator of an object
148293216
>>> f(a)
148293216
```

View or Shallow Copy

Different array objects can share the same data. The `view` method creates a new array object that looks at the same data.

```
>>> c = a.view()
>>> c is a
False
>>> c.base is a                          # c is viewing of the
data owned by a
True
>>> c.flags.owndata
False
>>>
>>> c.shape = 2,6                        # a's shape doesn't
change
>>> a.shape
(3, 4)
>>> c[0,4] = 1234                        # a's data changes
>>> a
array([[   0,    1,    2,    3],
       [1234,    5,    6,    7],
       [   8,    9,   10,   11]])
```

Slicing an array returns a view of it:

```
>>> s = a[:,1:3]
>>> s[:] = 10                            # s[:] is a view of s.
Note the difference between s=10 and s[:]=10
>>> a
array([[   0,   10,   10,    3],
       [1234,   10,   10,    7],
       [   8,   10,   10,   11]])
```

Deep Copy

The `copy` method makes a complete copy of the array and its data.

```
>>> d = a.copy()                         # a new array object
with new data is created
>>> d is a
False
>>> d.base is a                          # d doesn't share
anything with a
False
>>> d[0,0] = 9999
>>> a
array([[   0,   10,   10,    3],
       [1234,   10,   10,    7],
       [   8,   10,   10,   11]])
```

```
[    8,   10,   10,   11]])
```

Functions and Methods Overview

Here is a list of NumPy functions and methods names ordered in some categories. The names link to the Numpy Example List so that you can see the functions in action.

Array Creation
>   arange, array, copy, empty, empty like, eye, fromfile, fromfunction, identity, linspace, logspace, mgrid, ogrid, ones, ones like, r , zeros, zeros like

Conversions
>   astype, atleast 1d, atleast 2d, atleast 3d, mat

Manipulations
>   array split, column stack, concatenate, diagonal, dsplit, dstack, hsplit, hstack, item, newaxis, ravel, repeat, reshape, resize, squeeze, swapaxes, take, transpose, vsplit, vstack

Questions
>   all, any, nonzero, where

Ordering
>   argmax, argmin, argsort, max, min, ptp, searchsorted, sort

Operations
>   choose, compress, cumprod, cumsum, inner, fill, imag, prod, put, putmask, real, sum

Basic Statistics
>   cov, mean, std, var

Basic Linear Algebra
>   cross, dot, outer, svd, vdot

Less Basic

Broadcasting rules

Broadcasting allows universal functions to deal in a meaningful way with inputs that do not have exactly the same shape.

The first rule of broadcasting is that if all input arrays do not have the same number of dimensions, then a "1" will be repeatedly pre-pended to the shapes of the smaller arrays until all the arrays have the same number of dimensions.

The second rule of broadcasting ensures that arrays with a size of 1 along a particular dimension act as if they had the size of the array with the largest shape along that dimension. The value of the array element is assumed to be the same along that dimension for the "broadcasted" array.

After application of the broadcasting rules, the sizes of all arrays must match. More details can be found in this documentation.

Fancy indexing and index tricks

NumPy offers more indexing facilities than regular Python sequences. Basically, arrays can be indexed with integers and slices (as we saw before) and also with arrays of

integers and arrays of booleans.

Indexing with Arrays of Indices

```
>>> a = arange(12)**2                    # the first 12
square numbers
>>> i = array( [ 1,1,3,8,5 ] )           # an array of
indices
>>> a[i]                                  # the elements of a
at the positions i
array([ 1,  1,  9, 64, 25])
>>>
>>> j = array( [ [ 3, 4], [ 9, 7 ] ] )   # a bidimensional
array of indices
>>> a[j]                                  # the same shape as
j
array([[ 9, 16],
       [81, 49]])
```

When the indexed array `a` is multidimensional, a single array of indices reffers to the first dimension of `a`. The following example shows this behaviour by converting an image of labels into a color image using a palette.

```
>>> palette = array( [ [0,0,0],          # black
...                     [255,0,0],       # red
...                     [0,255,0],       # green
...                     [0,0,255],       # blue
...                     [255,255,255] ] ) # white
>>> image = array( [ [ 0, 1, 2, 0 ],     # each value
corresponds to a color in the palette
...                  [ 0, 3, 4, 0 ]  ] )
>>> palette[image]                        # the (2,4,3) color
image
array([[[  0,   0,   0],
        [255,   0,   0],
        [  0, 255,   0],
        [  0,   0,   0]],
       [[  0,   0,   0],
        [  0,   0, 255],
        [255, 255, 255],
        [  0,   0,   0]]])
```

We can also give indexes for more than one dimension. The arrays of indices for each dimension must have the same shape.

```
>>> a = arange(12).reshape(3,4)
>>> a
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>> i = array( [ [0,1],                   # indices for the
first dim of a
```

```
first dim of a
...                    [1,2] ] )
>>> j = array( [ [2,1],                        # indices for the second dim
...                    [3,3] ] )
>>>
>>> a[i,j]                                      # i and j must have equal shape
array([[ 2,  5],
       [ 7, 11]])
>>>
>>> a[i,2]
array([[ 2,  6],
       [ 6, 10]])
>>>
>>> a[:,j]
array([[[ 2,  1],
        [ 3,  3]],
       [[ 6,  5],
        [ 7,  7]],
       [[10,  9],
        [11, 11]]])
```

Naturally, we can put `i` and `j` in a sequence (say a list) and then do the indexing with the list.

```
>>> l = [i,j]
>>> a[l]                                        # equivalent to a[i,j]
array([[ 2,  5],
       [ 7, 11]])
```

However, we can not do this by putting `i` and `j` into an array, because this array will be interpreted as indexing the first dimension of a.

```
>>> s = array( [i,j] )
>>> a[s]                                        # not what we want
Traceback (most recent call last):
  File "<stdin>", line 1, in ?
IndexError: index (3) out of range (0<=index<=2) in dimension 0
>>>
>>> a[tuple(s)]                                 # same as a[i,j]
array([[ 2,  5],
       [ 7, 11]])
```

Another common use of indexing with arrays is the search of the maximum value of time-dependant series :

```
>>> time = linspace(20, 145, 5)                # time scale
>>> data = sin(arange(20)).reshape(5,4)        # 4 time-dependant series
```

```
>>> time
array([  20.  ,    51.25,    82.5 ,   113.75,   145.  ])
>>> data
array([[ 0.        ,   0.84147098,   0.90929743,   0.14112001],
       [-0.7568025 ,  -0.95892427,  -0.2794155 ,   0.6569866 ],
       [ 0.98935825,   0.41211849,  -0.54402111,  -0.99999021],
       [-0.53657292,   0.42016704,   0.99060736,   0.65028784],
       [-0.28790332,  -0.96139749,  -0.75098725,   0.14987721]])
>>>
>>> ind = data.argmax(axis=0)                        # index of the
maxima, for each series
>>> ind
array([2, 0, 3, 1])
>>>
>>> time_max = time[ ind]                             # time
corresponding to the maxima
>>>
>>> data_max = data[ind, xrange(data.shape[1])] # =>
data[ind[0],0], data[ind[1],1]...
>>>
>>> time_max
array([  82.5 ,    20.  ,   113.75,    51.25])
>>> data_max
array([ 0.98935825,   0.84147098,   0.99060736,   0.6569866 ])
>>>
>>> all(data_max == data.max(axis=0))
True
```

You can also use indexing with arrays as a target to assign to:

```
>>> a = arange(5)
>>> a
array([0, 1, 2, 3, 4])
>>> a[[1,3,4]] = 0
>>> a
array([0, 0, 2, 0, 0])
```

However, when the list of indices contains repetitions, the assignment is done several times, leaving behind the last value:

```
>>> a = arange(5)
>>> a[[0,0,2]]=[1,2,3]
>>> a
array([2, 1, 3, 3, 4])
```

This is reasonable enough, but watch out if you want to use python's += construct, as it may not do what you expect:

```
>>> a = arange(5)
>>> a[[0,0,2]]+=1
>>> a
```

```
array([1, 1, 3, 3, 4])
```

Even though 0 occurs twice in the list of indices, the 0th element is only incremented once. This is because python requires "a+=1" to be more or less equivalent to "a=a+1".

Indexing with Boolean Arrays

When we index arrays with arrays of (integer) indices we are providing the list of indices to pick. With boolean indices the approach is different; we explicitly choose which items in the array we want and which ones we don't.

The most natural way one can think of for boolean indexing is to use boolean arrays that have *the same shape* as the original array:

```
>>> a = arange(12).reshape(3,4)
>>> b = a > 4
>>> b                                       # b is a boolean
with a's shape
array([[False, False, False, False],
       [False, True, True, True],
       [True, True, True, True]], dtype=bool)
>>> a[b]                                     # 1d array with the
selected elements
array([ 5,  6,  7,  8,  9, 10, 11])
```

This property can be very useful in assignments:

```
>>> a[b] = 0                                # All elements of
'a' higher than 4 become 0
>>> a
array([[0, 1, 2, 3],
       [4, 0, 0, 0],
       [0, 0, 0, 0]])
```

You can look at the Mandelbrot set example to see how to use boolean indexing to generate an image of the Mandelbrot set.

The second way of indexing with booleans is more similar to integer indexing; for each dimension of the array we give a 1D boolean array selecting the slices we want.

```
>>> a = arange(12).reshape(3,4)
>>> b1 = array([False,True,True])           # first dim
selection
>>> b2 = array([True,False,True,False])     # second dim
selection
>>>
>>> a[b1,:]                                  # selecting rows
array([[ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>>
>>> a[b1]                                    # same thing
```

```
array([[ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>>
>>> a[:,b2]                              # selecting columns
array([[ 0,  2],
       [ 4,  6],
       [ 8, 10]])
>>>
>>> a[b1,b2]                             # a weird thing to
do
array([ 4, 10])
```

Note that the length of the 1D boolean array has to coincide with the length of the dimension (or axis) you want to slice. In the previous example, b1 is a 1-rank array with length 3 (the number of *rows* in a), and b2 (of length 4) is suitable to index the 2nd rank (columns) of a.

The ix_() function

The ix_ function can be used to combine different vectors so as to obtain the result for each n-uplet. For example, if you want to compute all the a+b*c for all the triplets taken from each of the vectors a, b and c:

```
>>> a = array([2,3,4,5])
>>> b = array([8,5,4])
>>> c = array([5,4,6,8,3])
>>> ax,bx,cx = ix_(a,b,c)
>>> ax.shape, bx.shape, cx.shape
((4, 1, 1), (1, 3, 1), (1, 1, 5))
>>> result = ax+bx*cx
>>> result
array([[[42, 34, 50, 66, 26],
        [27, 22, 32, 42, 17],
        [22, 18, 26, 34, 14]],
       [[43, 35, 51, 67, 27],
        [28, 23, 33, 43, 18],
        [23, 19, 27, 35, 15]],
       [[44, 36, 52, 68, 28],
        [29, 24, 34, 44, 19],
        [24, 20, 28, 36, 16]],
       [[45, 37, 53, 69, 29],
        [30, 25, 35, 45, 20],
        [25, 21, 29, 37, 17]]])
>>> result[3,2,4]
17
>>> a[3]+b[2]*c[4]
17
```

You could also implement the reduce like that:

```
def ufunc_reduce(ufct, *vectors):
```

```
    vs = ix_(*vectors)
    r = ufct.identity
    for v in vs:
        r = ufct(r,v)
    return r
```

and then use it as:

```
>>> ufunc_reduce(add,a,b,c)
array([[[15, 14, 16, 18, 13],
        [12, 11, 13, 15, 10],
        [11, 10, 12, 14,  9]],
       [[16, 15, 17, 19, 14],
        [13, 12, 14, 16, 11],
        [12, 11, 13, 15, 10]],
       [[17, 16, 18, 20, 15],
        [14, 13, 15, 17, 12],
        [13, 12, 14, 16, 11]],
       [[18, 17, 19, 21, 16],
        [15, 14, 16, 18, 13],
        [14, 13, 15, 17, 12]]])
```

The big advantage of this version of reduce compared to the normal ufunc.reduce is that you don't need to create an argument array the size of the output time the number of vectors, but you make use of the Broadcasting Rules.

Indexing with strings

See RecordArrays.

Linear Algebra

Work in progress. Basic linear algebra to be included here.

Simple Array Operations

See linalg.py in numpy folder for more.

```
>>> from numpy import *
>>> from numpy.linalg import *

>>> a = array([[1.0, 2.0], [4.0, 3.0]])
>>> print a
[[ 1.  2.]
 [ 3.  4.]]

>>> a.transpose()
array([[ 1.,  3.],
       [ 2.,  4.]])

>>> inv(a)
```

```
>>> inv(a)
array([[-2. ,  1. ],
       [ 1.5, -0.5]])

>>> u = eye(2) # unit 2x2 matrix
>>> j = array([[0.0, -1.0], [1.0, 0.0]])

>>> dot (j, j) # matrix product
array([[-1.,  0.],
       [ 0., -1.]])

>>> trace(u)  # trace
2.0

>>> y = array([[5.], [7.]])
>>> solve(a, y)
array([[-3.],
       [ 4.]])

>>> eig(j)
(array([ 0.+1.j,  0.-1.j]),
array([[ 0.70710678+0.j,  0.70710678+0.j],
       [ 0.00000000-0.70710678j,  0.00000000+0.70710678j]]))
Parameters:
    square matrix

Returns
    The eigenvalues, each repeated according to its multiplicity.

    The normalized (unit "length") eigenvectors, such that the
    column ``v[:,i]`` is the eigenvector corresponding to the
    eigenvalue ``w[i]`` .
```

The Matrix Class

Here is a short intro to the Matrix class.

```
>>> A = matrix('1.0 2.0; 3.0 4.0')
>>> A
[[ 1.  2.]
 [ 3.  4.]]
>>> type(A)  # file where class is defined
<class 'numpy.matrixlib.defmatrix.matrix'>

>>> A.T  # transpose
[[ 1.  3.]
 [ 2.  4.]]

>>> X = matrix('5.0 7.0')
>>> Y = X.T
>>> Y
```

```
 [[5.]
  [7.]]

>>> print A*Y  # matrix multiplication
[[19.]
 [43.]]

>>> print A.I  # inverse
[[-2.    1. ]
 [ 1.5 -0.5]]

>>> solve(A, Y)  # solving linear equation
matrix([[-3.],
        [ 4.]])
```

Indexing: Comparing Matrices and 2-d Arrays

Note that there are some important differences between numpy arrays and matrices. NumPy provides two fundamental objects: an N-dimensional array object and a universal function object. Other objects are built on top of these. In particular, matrices are 2-dimensional array objects that inherit from the NumPy array object. For both arrays and matrices, indices must consist of a proper combination of one or more of the following: integer scalars, ellipses, a list of integers or boolean values, a tuple of integers or boolean values, and a 1-dimensional array of integer or boolean values. A matrix can be used as an index for matrices, but commonly an array, list, or other form is needed to accomplish a given task.

As usual in Python, indexing is zero-based. Traditionally we represent a 2-d array or matrix as a rectangular array of rows and columns, where movement along axis 0 is movement across rows, while movement along axis 1 is movement across columns.

Let's make an array and matrix to slice:

```
>>> A = arange(12)
>>> A
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
>>> A.shape = (3,4)
>>> M = mat(A.copy())
>>> print type(A),"  ",type(M)
<type 'numpy.ndarray'>    <class 'numpy.core.defmatrix.matrix'>
>>> print A
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
>>> print M
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
```

Now, let's take some simple slices. Basic slicing uses slice objects or integers. For example, the evaluation of A[:] and M[:] will appear familiar from Python indexing, however it is important to note that slicing numpy arrays does *not* make a copy of the

data: slicing provides a new view of the same data.

```
>>> print A[:]; print A[:].shape
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
(3, 4)
>>> print M[:]; print M[:].shape
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
(3, 4)
```

Now for something that differs from Python indexing: you may use comma-separated indices to index along multiple axes at the same time.

```
>>> print A[:,1]; print A[:,1].shape
[1 5 9]
(3,)
>>> print M[:,1]; print M[:,1].shape
[[1]
 [5]
 [9]]
(3, 1)
```

Notice the difference in the last two results. Use of a single colon for the 2-d array produces a 1-dimensional array, while for a matrix it produces a 2-dimensional matrix. A slice of a matrix will always produce a matrix. For example, a slice M[2,:] produces a matrix of shape (1,4). In contrast, a slice of an array will always produce an array of the lowest possible dimension. For example, if C were a three dimensional array, C[...,1] produces a 2-d array while C[1,:,1] produces a 1 dimensional array. From this point on, we will show results only for the array slice if the results for the corresponding matrix slice are identical.

Lets say that we wanted the 1st and 3rd column of an array. One way is to slice using a list:

```
>>> A[:,[1,3]]
array([[ 1,  3],
       [ 5,  7],
       [ 9, 11]])
```

A slightly more complicated way is to use the take() method:

```
>>> A[:,].take([1,3],axis=1)
array([[ 1,  3],
       [ 5,  7],
       [ 9, 11]])
```

If we wanted to skip the first row, we could use:

```
>>> A[1:,].take([1,3],axis=1)
array([[  5,   7],
       [  9,  11]])
```

Or we could simply use A[1:,[1,3]]. Yet another way to slice the above is to use a cross product:

```
>>> A[ix_((1,2),(1,3))]
array([[  5,   7],
       [  9,  11]])
```

For the reader's convenience, here is our array again:

```
>>> print A
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
```

Now let's do something a bit more complicated. Lets say that we want to retain all columns where the first row is greater than 1. One way is to create a boolean index:

```
>>> A[0,:]>1
array([False, False, True, True], dtype=bool)
>>> A[:,A[0,:]>1]
array([[ 2,   3],
       [ 6,   7],
       [10,  11]])
```

Just what we wanted! But indexing the matrix is not so convenient.

```
>>> M[0,:]>1
matrix([[False, False, True, True]], dtype=bool)
>>> M[:,M[0,:]>1]
matrix([[2, 3]])
```

The problem of course is that slicing the matrix slice produced a matrix. But matrices have a convenient 'A' attribute whose value is the array representation, so we can just do this instead:

```
>>> M[:,M.A[0,:]>1]
matrix([[ 2,   3],
        [ 6,   7],
        [10,  11]])
```

If we wanted to conditionally slice the matrix in two directions, we must adjust our strategy slightly. Instead of

```
>>> A[A[:,0]>2,A[0,:]>1]
array([ 6, 11])
>>> M[M.A[:,0]>2,M.A[0,:]>1]
```

```
>>> M[M.A[:,0]>2,M.A[0,:]>1]
matrix([[ 6, 11]])
```

we need to use the cross product 'ix_':

```
>>> A[numpy.ix_(A[:,0]>2,A[0,:]>1)]
array([[ 6,  7],
       [10, 11]])
>>> M[numpy.ix_(M.A[:,0]>2,M.A[0,:]>1)]
matrix([[ 6,  7],
        [10, 11]])
```

Tricks and Tips

Here we give a list of short and useful tips.

"Automatic" Reshaping

To change the dimensions of an array you can omit one of the sizes which will then be deduced automatically:

```
>>> a = arange(30)
>>> a.shape = 2,-1,3  # -1 means "whatever is needed"
>>> a.shape
(2, 5, 3)
>>> a
array([[[ 0,  1,  2],
        [ 3,  4,  5],
        [ 6,  7,  8],
        [ 9, 10, 11],
        [12, 13, 14]],
       [[15, 16, 17],
        [18, 19, 20],
        [21, 22, 23],
        [24, 25, 26],
        [27, 28, 29]]])
```

Vector Stacking

How to construct a 2D array from a list of equally-sized row vectors? In matlab this is quite easy: if x and y are two vectors of the same length you only need do m=[x;y]. In NumPy this works via the functions column_stack, dstack, hstack and vstack, depending on which dimension the stacking is to be done. For example:

```
x = arange(0,10,2)          # x=([0,2,4,6,8])
y = arange(5)               # y=([0,1,2,3,4])
m = vstack([x,y])           # m=([[0,2,4,6,8],
                            #     [0,1,2,3,4]])
xy = hstack([x,y])          # xy =
([0,2,4,6,8,0,1,2,3,4])
```

```
\[ . ,  ,  ,  ,  ,  , , , , ])
```

The logic behind those functions in more than two dimensions can be strange.

See also NumPy for Matlab Users and add there your new findings 😊

Histograms

The NumPy's function `histogram` applied to an array returns a pair of vectors: the histogram of the array and the vector of bins. Beware: `matplotlib` has also a function to build histograms (called `hist`, as in matlab) different from that of NumPy. The main difference is that `pylab.hist` plots the histogram automatically, while `numpy.histogram` only generates the data.

```python
import numpy
import pylab
# Build a vector of 10000 normal deviates with variance 0.5^2 and
mean 2
mu, sigma = 2, 0.5
v = numpy.random.normal(mu,sigma,10000)
# Plot a normalized histogram with 50 bins
pylab.hist(v, bins=50, normed=1)       # matplotlib version (plot)
pylab.show()
# Compute the histogram with numpy and then plot it
(n, bins) = numpy.histogram(v, bins=50, normed=True)  # NumPy
version (no plot)
pylab.plot(.5*(bins[1:]+bins[:-1]), n)
pylab.show()
```

References

- The Python tutorial.
- The Numpy Example List.
- The inexistent NumPy Tutorial at scipy.org, where we can find the old Numeric documentation.
- The Guide to NumPy book.
- The SciPy Tutorial and a SciPy course online
- NumPy for Matlab Users.
- A matlab, R, IDL, NumPy/SciPy dictionary.

Tentative NumPy Tutorial (last edited 2011-05-02 04:36:55 by EdwardGrutman)